

Scalable Monte Carlo

Babak Shahbaba

Department of Statistics, UCI

May 4, 2017

Introduction

Bayesian inference

- In Bayesian statistics we make inference based on posterior probability distribution $P(\theta|y)$:

$$\begin{aligned} P(\theta|y) &= \frac{P(y|\theta)P(\theta)}{P(y)} \\ &\propto P(y|\theta)P(\theta) \end{aligned}$$

- For example, we can predict future observations, \tilde{y} , given the observed data y :

$$P(\tilde{y}|y) = \int_{\theta} P(\tilde{y}|\theta)P(\theta|y)d\theta$$

- **Main challenge:** inference almost always involves intractable integrals

- In Bayesian statistics we make inference based on posterior probability distribution $P(\theta|y)$:

$$\begin{aligned} P(\theta|y) &= \frac{P(y|\theta)P(\theta)}{P(y)} \\ &\propto P(y|\theta)P(\theta) \end{aligned}$$

- For example, we can predict future observations, \tilde{y} , given the observed data y :

$$P(\tilde{y}|y) = \int_{\theta} P(\tilde{y}|\theta)P(\theta|y)d\theta$$

- **Main challenge:** inference almost always involves intractable integrals

Monte Carlo approximation

- In general, we can approximate

$$\mu = \int_{\mathcal{X}} h(x)f(x)dx$$

using iid samples $x^{(1)}, x^{(2)}, \dots, x^{(m)}$ from the distribution with density $f(x)$:

$$\hat{\mu} = \frac{1}{m}[h(x^{(1)}) + h(x^{(2)}) + \dots + h(x^{(m)})]$$

- **Main challenge:** usually we cannot generate iid samples from the distribution

Monte Carlo approximation

- In general, we can approximate

$$\mu = \int_{\mathcal{X}} h(x)f(x)dx$$

using iid samples $x^{(1)}, x^{(2)}, \dots, x^{(m)}$ from the distribution with density $f(x)$:

$$\hat{\mu} = \frac{1}{m}[h(x^{(1)}) + h(x^{(2)}) + \dots + h(x^{(m)})]$$

- **Main challenge:** usually we cannot generate iid samples from the distribution

Markov chain Monte Carlo

- We use Markov chains to generate samples from the distribution

Markov Chains



Markov Chain Monte Carlo



- **Main challenge:** finding a good transition probability


Markov chain Monte Carlo

- We use Markov chains to generate samples from the distribution

Markov Chains



Markov Chain Monte Carlo



- **Main challenge:** finding a good transition probability

The Metropolis Algorithm

- Specify a symmetric transition probability $g(\theta, \theta^*)$ and repeat the following steps for many iterations:

- Given our current state θ , we propose a new state θ^* according to g .
- Calculated the acceptance probability,

$$\begin{aligned}a(\theta, \theta^*) &= \min\left(1, \frac{f(\theta^*)}{f(\theta)}\right) \\ &= \min\{1, \exp(\log[f(\theta^*)] - \log[f(\theta)])\}\end{aligned}$$

- Accept the proposed state θ^* as the new state with probability $a(\theta, \theta^*)$ or remain at state θ .
- For asymmetrical proposal distribution, we use Metropolis-Hastings (MH),

$$a(\theta, \theta^*) = \min\left(1, \frac{f(\theta^*)g(\theta^*, \theta)}{f(\theta)g(\theta, \theta^*)}\right)$$

- Main challenge:** finding the right proposal-generating mechanism, $g(\theta, \theta^*)$

The Metropolis Algorithm

- Specify a symmetric transition probability $g(\theta, \theta^*)$ and repeat the following steps for many iterations:

- Given our current state θ , we propose a new state θ^* according to g .

- Calculated the acceptance probability,

$$\begin{aligned}a(\theta, \theta^*) &= \min\left(1, \frac{f(\theta^*)}{f(\theta)}\right) \\ &= \min\{1, \exp(\log[f(\theta^*)] - \log[f(\theta)])\}\end{aligned}$$

- Accept the proposed state θ^* as the new state with probability $a(\theta, \theta^*)$ or remain at state θ .


- For asymmetrical proposal distribution, we use Metropolis-Hastings (MH),

$$a(\theta, \theta^*) = \min\left(1, \frac{f(\theta^*)g(\theta^*, \theta)}{f(\theta)g(\theta, \theta^*)}\right)$$

- Main challenge:** finding the right proposal-generating mechanism, $g(\theta, \theta^*)$

Random walk Metropolis



- A simple proposal generating mechanism is random walk: $\theta^* \sim N(\theta, \epsilon^2 I)$



Hamiltonian Monte Carlo (HMC)


Hamiltonian Monte Carlo

- HMC proposes states that are distant from the current state, but nevertheless have a high probability of acceptance.



Hamiltonian Monte Carlo


- The sampler is viewed as a dynamic system moving on a surface defined by the *energy* function U : negative log density of the target distribution



- Distant proposals are found by numerically simulating Hamiltonian dynamics for some specified amount of fictitious time

Hamiltonian Monte Carlo

- The sampler is viewed as a dynamic system moving on a surface defined by the *energy* function U : negative log density of the target distribution



- Distant proposals are found by numerically simulating Hamiltonian dynamics for some specified amount of fictitious time

Posterior sampling

- For Bayesian inference, posterior distribution is the target distribution

Potential energy

$$U(\theta) = - \sum_{i=1}^N \log P(y_i|\theta) - \log P(\theta)$$

- We augment the parameter space with fictitious momentum variables

Kinetic energy

$$K(p) = \frac{1}{2} p^\top M^{-1} p$$

- Define the Hamiltonian function $H(\theta, p) = U(\theta) + K(p)$
- The joint density of (θ, p) is

$$P(\theta, p) \propto \exp\{-H(\theta, p)\} = \exp\{-U(\theta)\} \cdot \exp\{-K(p)\}$$

- The marginal distribution of θ is the posterior distribution

Hamiltonian Dynamics

- We can generate a proposal by starting from the current state at time 0 and moving to the state at time t :

$$(\theta, p) = (\theta^{(0)}, p^{(0)}) \xrightarrow{\text{HD}} (\theta^{(t)}, p^{(t)}) = (\theta^*, p^*)$$

- *Hamilton's equations* determine how θ and p change over [fictitious] time

Hamilton's equations

$$\begin{aligned}\frac{d\theta_j}{dt} &= \frac{\partial H}{\partial p_j} = [M^{-1}p]_j \\ \frac{dp_j}{dt} &= -\frac{\partial H}{\partial \theta_j} = -\frac{\partial U}{\partial \theta_j}\end{aligned}$$

- **Important properties:**
 - ▶ **Reversibility:** the target distribution remain invariant
 - ▶ **Volume preservation:** the Jacobin determinant is 1
 - ▶ **Conservation of Hamiltonian:** the acceptance rate is one; θ^* is the next sample if HD is analytically solvable

- Numerical integration is employed when analytic solution is not available

Leapfrog

$$p_j(t + \epsilon/2) = p_j(t) - (\epsilon/2) \frac{\partial U}{\partial \theta_j}(\theta(t))$$

$$\theta_j(t + \epsilon) = \theta_j(t) + \epsilon \frac{\partial K}{\partial p_j}(p(t + \epsilon/2))$$

$$p_j(t + \epsilon) = p_j(t + \epsilon/2) - (\epsilon/2) \frac{\partial U}{\partial \theta_j}(\theta(t + \epsilon))$$

- **Important properties:**

- ▶ **Stability:** numerically stable if ϵ is appropriately chosen
- ▶ **Reversibility and Volume preservation:** still hold
- ▶ **Conservation of Hamiltonian:** broken, but can be corrected by MH correction step with acceptance rate

$$\alpha = \min[1, \exp(-H(\theta^*, p^*) + H(\theta, p))]$$

Algorithm 1: HMC algorithm

Initialize $\theta^{(0)} = \text{current } \theta$

Sample new momentum $p^{(0)} \sim \mathcal{N}(0, M = I)$

Calculate current $H(\theta^{(0)}, p^{(0)}) = U(\theta^{(0)}) + \frac{1}{2}(p^{(0)})^\top p^{(0)}$

for $\ell = 1$ to L (leapfrog steps) **do**

$$p^{(\ell+\frac{1}{2})} = p^{(\ell)} - \epsilon/2 \nabla_{\theta} U(\theta^{(\ell)})$$

$$\theta^{(\ell+1)} = \theta^{(\ell)} + \epsilon p^{(\ell+\frac{1}{2})}$$

$$p^{(\ell+1)} = p^{(\ell+\frac{1}{2})} - \epsilon/2 \nabla_{\theta} U(\theta^{(\ell+1)})$$

end for

Accept or reject according to the Metropolis acceptance probability

Example: logistic regression with $N(0, \sigma^2 I)$ prior

$$\nabla_{\beta_j} U(\beta) = - \sum_{i=1}^N [y_i - \frac{\exp(x_i \beta)}{1 + \exp(x_i \beta)}] x_{ij} + \beta_j / \sigma^2$$

A special case: Langevin Monte Carlo

- A special case: $L = 1$ and $M = I$
- This is called Langevin Monte Carlo,

Langevin dynamics

$$\begin{aligned}\theta^* &= \theta - \frac{\epsilon^2}{2} \nabla_{\theta} U(\theta) + \epsilon p \\ p^* &= p - \frac{\epsilon}{2} \nabla_{\theta} U(\theta) - \frac{\epsilon}{2} \nabla_{\theta} U(\theta^*)\end{aligned}$$

- Alternatively, we could ignore the momentum variable p and use the following asymmetrical proposal with MH acceptance probability

$$\theta^* \sim N\left(\theta - \frac{\epsilon^2}{2} \nabla_{\theta} U(\theta), \epsilon^2 I\right)$$

- Dropping the accept/reject step leads to an approximate Langevin method (see Neal, 1993)




A more general case: Riemannian Manifold HMC

- Girolami and Calderhead (2011) have introduced a new method, called Riemannian Manifold HMC (RMHMC)
- They argue that it is more natural to put the Hamiltonian dynamic on Riemannian manifold of distributions rather than Euclidean space
- They follow Amari (2000) and use the Fisher information matrix, $G(\theta) = E[\nabla_{\theta}^2 U(\theta)]$, as a metric on the manifold
- That is, they use position specific mass matrix, $M = G(\theta)$
- Example: logistic regression

$$G_{jk}(\beta) = \sum_{i=1}^N x_{ij} x_{ik} \frac{\exp(x_i \beta)}{[1 + \exp(x_i \beta)]^2}, \quad j \neq k$$

- We can explore the parameter space more efficiently by exploiting its geometric properties
- The resulting dynamics is non-separable so instead of the standard leapfrog method we need to use the *generalized* leapfrog method

HMC vs. RMHMC



A main challenge: high computational cost

- For high-dimensional problems (big n and/or big d) and complex models, these methods tend to be computationally expensive
- We have proposed several variations of HMC:
 - ▶ Split HMC (S. et al., 2011)
 - ▶ Lagrangian Monte Carlo (Lan, et al., 2012)
 - ▶ Spherical HMC (Lan et al., 2013)
 - ▶ Wormhole HMC (Lan et al., 2013)
 - ▶ HMC with precomputing strategy (Zhang et al., 2015)
 - ▶ HMC with surrogate functions (Zhang et al., 2015)

Scalable HMC

Subsampling

- In recent years, computational methods based on mini-batches of data have been quite successful
 - ▶ The underlying assumption: there is redundancy in big data
 - ▶ The overall information can be retrieved from a small subset
 - ▶ We can approximate functions at low computational cost
- Welling and Teh (2011) used this approach (stochastic gradient) for Langevin dynamics using mini-batches of size n from N observations

$$\theta^* = \theta + \frac{\epsilon^2}{2} (\nabla_{\theta} P(\theta) + \frac{N}{n} \sum_{i=1}^n \nabla_{\theta} \log P(x_i | \theta)) + \epsilon p$$

- They also dropped the accept/reject step

- In recent years, computational methods based on mini-batches of data have been quite successful
 - ▶ The underlying assumption: there is redundancy in big data
 - ▶ The overall information can be retrieved from a small subset
 - ▶ We can approximate functions at low computational cost
- Welling and Teh (2011) used this approach (stochastic gradient) for Langevin dynamics using mini-batches of size n from N observations

$$\theta^* = \theta + \frac{\epsilon^2}{2} (\nabla_{\theta} P(\theta) + \frac{N}{n} \sum_{i=1}^n \nabla_{\theta} \log P(x_i | \theta)) + \epsilon p$$

- They also dropped the accept/reject step

- In recent years, computational methods based on mini-batches of data have been quite successful
 - ▶ The underlying assumption: there is redundancy in big data
 - ▶ The overall information can be retrieved from a small subset
 - ▶ We can approximate functions at low computational cost
- Welling and Teh (2011) used this approach (stochastic gradient) for Langevin dynamics using mini-batches of size n from N observations

$$\theta^* = \theta + \frac{\epsilon^2}{2} (\nabla_{\theta} P(\theta) + \frac{N}{n} \sum_{i=1}^n \nabla_{\theta} \log P(x_i | \theta)) + \epsilon p$$

- They also dropped the accept/reject step

Focusing on parameter space

- Finding optimum subsets by exploiting regularity in data space is difficult
- Using random subsets could lead to non-ignorable loss of information
- Therefore, we previously proposed to identify a subset of influential points to split the Hamiltonian function (Leimkuler and Reich, 2004) into two parts (S. et al., 2011)
- Recently, we have switched our focus from data space to parameter space
 - ▶ We exploit the smoothness and regularity of parameter space
 - ▶ We precompute functions (e.g., gradient) on a relatively small sample of parameters
 - ▶ MCMC algorithms use these precomputed values to approximate functions
 - ▶ We use the exact target distribution for the accept/reject step to ensure convergence to the right stationary distribution

Focusing on parameter space

- Finding optimum subsets by exploiting regularity in data space is difficult
- Using random subsets could lead to non-ignorable loss of information
- Therefore, we previously proposed to identify a subset of influential points to split the Hamiltonian function (Leimkuler and Reich, 2004) into two parts (S. et al., 2011)
- Recently, we have switched our focus from data space to parameter space
 - ▶ We exploit the smoothness and regularity of parameter space
 - ▶ We precompute functions (e.g., gradient) on a relatively small sample of parameters
 - ▶ MCMC algorithms use these precomputed values to approximate functions
 - ▶ We use the exact target distribution for the accept/reject step to ensure convergence to the right stationary distribution

Focusing on parameter space

- Finding optimum subsets by exploiting regularity in data space is difficult
- Using random subsets could lead to non-ignorable loss of information
- Therefore, we previously proposed to identify a subset of influential points to split the Hamiltonian function (Leimkuler and Reich, 2004) into two parts (S. et al., 2011)
- Recently, we have switched our focus from data space to parameter space
 - ▶ We exploit the smoothness and regularity of parameter space
 - ▶ We precompute functions (e.g., gradient) on a relatively small sample of parameters
 - ▶ MCMC algorithms use these precomputed values to approximate functions
 - ▶ We use the exact target distribution for the accept/reject step to ensure convergence to the right stationary distribution

Focusing on parameter space

- Finding optimum subsets by exploiting regularity in data space is difficult
- Using random subsets could lead to non-ignorable loss of information
- Therefore, we previously proposed to identify a subset of influential points to split the Hamiltonian function (Leimkuler and Reich, 2004) into two parts (S. et al., 2011)
- Recently, we have switched our focus from data space to parameter space
 - ▶ We exploit the smoothness and regularity of parameter space
 - ▶ We precompute functions (e.g., gradient) on a relatively small sample of parameters
 - ▶ MCMC algorithms use these precomputed values to approximate functions
 - ▶ We use the exact target distribution for the accept/reject step to ensure convergence to the right stationary distribution

Naive Grid Approximation (GHMC)

$$\begin{aligned}\frac{dp_j}{dt} &= -\frac{\partial U}{\partial \theta_j} \\ \frac{d\theta_j}{dt} &= [M^{-1}p]_j\end{aligned}$$

Denote

Force


$$F = -\nabla U$$

- piecewise constant approximation

$$\tilde{F}(\theta) = F_{i,j} \triangleq F(c_{i,j}), \quad \text{if } \theta \in C_{i,j}$$


- piecewise linear approximation

$$\tilde{F}(\theta) = F_{i,j} + \nabla F_{i,j} \cdot (q - c_{i,j}), \quad \text{if } \theta \in C_{i,j}$$




Force map of a logistic regression model

True $\beta = (-1, 1)$




HMC




GHMC

Logistic regression



HMC




GHMC

Banana-shaped distribution

| Experiment | Method | AR | s/Iteration | min(ESS)/s | Sped-up |
|------------|--------|--------|-------------|------------|------------|
| LR | HMC | 0.9225 | 7.0157E-4 | 1425.4 | 1 |
| | GHMC | 0.7981 | 3.318E-4 | 3013.9 | 2.1 |
| BD | HMC | 0.9353 | 3.8703E-4 | 962.1 | 1 |
| | GHMC | 0.6587 | 1.4498E-4 | 1651.6 | 1.7 |


Sparse Grid for Higher Dimensions (SGHMC)

- In higher dimensions, we use a sparse grid interpolation method based on *Smolyak's* formula (Barthelmann, 2000; Bungartz, 1998 & 2004)
- It employs $\mathcal{O}(N \cdot (\log(N))^{d-1})$ points only, with approximation accuracy preserved up to a logarithmic factor




HMC

Gaussian Process model




SGHMC



HMC

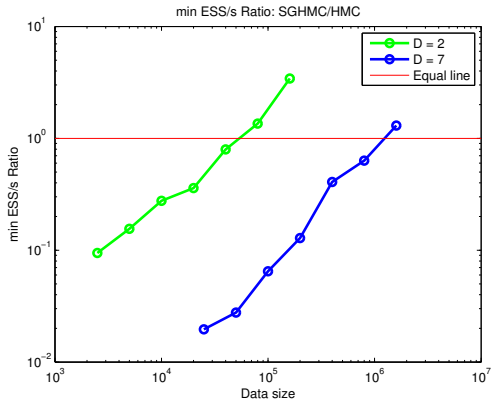
Elliptic PDE inverse problem



SGHMC

| Experiment | Method | AR | s/Iteration | min(ESS)/s | Speed-up |
|------------|--------|--------|-------------|------------|------------|
| GP | HMC | 0.9472 | 2.3547E-1 | 1.3 | 1 |
| | SGHMC | 0.7066 | 2.9851E-2 | 8.7 | 6.7 |
| ePDE | HMC | 0.7719 | 2.02E-1 | 1.5 | 1 |
| | SGHMC | 0.6141 | 6.1952E-2 | 4.3 | 2.9 |

- However, this approach reaches its limit quite fast



HMC with Surrogate Functions

- In recent years, several methods have been proposed based on constructing *surrogate* Hamiltonians using Gaussian process models (Rasmussen, 2003; Meeds and Welling, 2015; Lan et. al., 2015)
- We have instead used a simple generalized additive model, which can be regarded as a shallow neural network,

$$\tilde{U}(\theta) = \sum_{i=1}^s v_i g(\mathbf{w}_i \cdot \theta + d_i) + d_0$$

with the softplus function: $g(z) = \log(1 + \exp(z))$

Extreme Learning Machine (ELM)

For training, we randomly assign input weights and biases, and then obtain the least-square estimates of the output weights v

ELM (Huang, 2006)

Given a training set $\mathcal{T} = \{(I_j, t_j) | I_j \in \mathbb{R}^n, t_j \in \mathbb{R}^m, j = 1, \dots, N\}$, activation function $\sigma(x)$ and hidden node number s


- 1 Randomly assign input weight w_i and bias d_i , $i = 1, \dots, s$
- 2 Calculate the hidden layer output matrix H

$$H_{ji} = \sigma(w_i I_j + d_i), \quad i = 1, \dots, s, j = 1, \dots, N$$


- 3 Calculate the output weight v

$$v = H^\dagger T, \quad T = [t_1, t_2, \dots, t_N]^T$$

where H^\dagger is the *Moore-Penrose generalized inverse* of matrix H



Target function






Neural network approximation

- The training process (using pre-convergence samples) and the approximation of functions in the sampling phase can be easily incorporated in HMC
- The approximate geometric information (e.g, gradient and Hessian) is obtained by differentiating the neural network directly,

$$\frac{\partial \tilde{U}}{\partial \theta} = \sum_{i=1}^s v_i g'(\mathbf{w}_i \cdot \boldsymbol{\theta} + d_i) \mathbf{w}_i$$

- Easy generalization to Riemannian Manifold HMC (NNS-RMHMC)

Surrogate Induced Hamiltonian Flow



Experiments

| Experiment | Method | AP | s/lter | min(ESS)/s | Spdd-up |
|---------------------|-----------|--------|-----------|------------|--------------|
| LR (Simulation) | HMC | 0.6656 | 3.573E-01 | 1.45 | 1 |
| | RMHMC | 0.8032 | 3.794 | 0.06 | 0.04 |
| | NNS-HMC | 0.6726 | 1.364E-02 | 37.83 | 26.09 |
| | NNS-RMHMC | 0.8162 | 1.027E-01 | 2.17 | 1.50 |
| LR (Bank Marketing) | HMC | 0.8038 | 7.400E-02 | 6.51 | 1 |
| | RMHMC | 0.9210 | 6.305E-01 | 0.56 | 0.08 |
| | NNS-HMC | 0.7944 | 7.508E-03 | 58.22 | 8.94 |
| | NNS-RMHMC | 0.9064 | 2.741E-02 | 14.41 | 2.21 |
| LR (Adult Data) | HMC | 0.8300 | 7.898E-02 | 0.21 | 1 |
| | RMHMC | 0.8526 | 5.842E-01 | 1.06 | 4.81 |
| | NNS-HMC | 0.8096 | 9.914E-03 | 2.66 | 12.09 |
| | NNS-RMHMC | 0.8400 | 3.300E-02 | 18.68 | 84.90 |
| Elliptic PDE | HMC | 0.7077 | 1.568 | 0.061 | 1 |
| | RMHMC | 0.8014 | 4.388 | 0.228 | 3.74 |
| | NNS-HMC | 0.7138 | 7.419E-02 | 1.410 | 23.11 |
| | NNS-RMHMC | 0.6584 | 9.720E-02 | 4.375 | 71.72 |

Variational HMC

- Alternatively, we can make inference based on an approximate distribution, similar to variational Bayes, but with a better and more flexible approximation (see for example, de Freitas et al., 2001; Salimans et al., 2015)
- For variational Bayes, we typically use a parametrized distribution $q_{\eta}(\theta)$ to approximate the target posterior $p(\theta|Y)$ by minimizing the KL divergence
- Here, we use the approximate distribution based on our neural network model

$$Q_v(\theta) \propto \exp(-\tilde{U}(\theta)) = \exp\left[-\sum_{i=1}^s v_i g(\mathbf{w}_i \cdot \theta + d_i) + \phi(v)\right]$$

- This is simply a flexible exponential family model

Free-form variational Bayes

- To find Q_v , we follow Hyvarinen (2005) and minimize the score-matching distance

$$\tilde{D}_{SM}(P(\theta|Y)||Q_v(\theta)) = \frac{1}{2} \int Q_v(\theta) \|\nabla_{\theta} \tilde{U}(\theta) - \nabla_{\theta} U(\theta)\|^2 d\theta$$

- For this, we use HMC to generate samples from Q_v

$$\begin{aligned} \frac{d\theta}{dt} &= \frac{\partial \tilde{H}}{\partial p} = M^{-1}p \\ \frac{dp}{dt} &= -\frac{\partial \tilde{H}}{\partial \theta} = -\nabla_{\theta} \tilde{U}(\theta) \end{aligned}$$

where the modified Hamiltonian is

$$\tilde{H}(\theta, p) = \tilde{U}(\theta) + K(p)$$

- Then minimize the regularized empirical distance

$$\hat{v} = \arg \min_v \frac{1}{2} \sum_{n=1}^t \|\nabla_{\theta} \tilde{U}(\theta_n) - \nabla_{\theta} U(\theta_n)\|^2 + \frac{\lambda}{2} \|v\|^2$$

Online updating of the weight vector

- Given the current weight vector $v^{(t)}$ and a new training data point $(\theta_{t+1}, \nabla_{\theta} U(\theta_{t+1}))$, the updating formula for the estimator is given by

$$v^{(t+1)} = v^{(t)} + W^{(t+1)}(\nabla_{\theta} U(\theta_{t+1}) - A_{t+1}v^{(t)})$$

where

$$W^{(t+1)} = C^{(t)} A'_{t+1} \left[I_d + A_{t+1} C^{(t)} A'_{t+1} \right]^{-1}$$

$$A_{t+1} = (A_1(\theta_{t+1}), \dots, A_s(\theta_{t+1}))$$

with $A_i(\theta_{t+1}) := \sigma'(w_i \cdot \theta_{t+1} + d_i) w_i$, and $C^{(t)}$ can be updated by *Sherman-Morrison-Woodbury* formula:

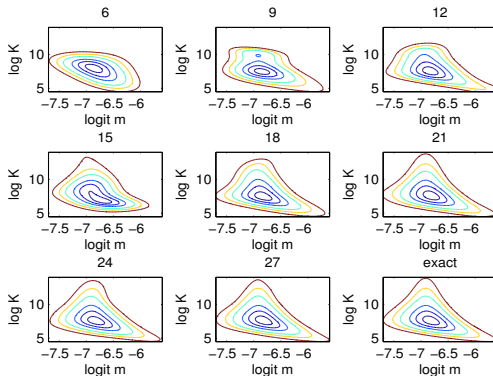
$$C^{(t+1)} = C^{(t)} - W^{(t+1)} A_{t+1} C^{(t)}$$

Example: a beta-binomial model

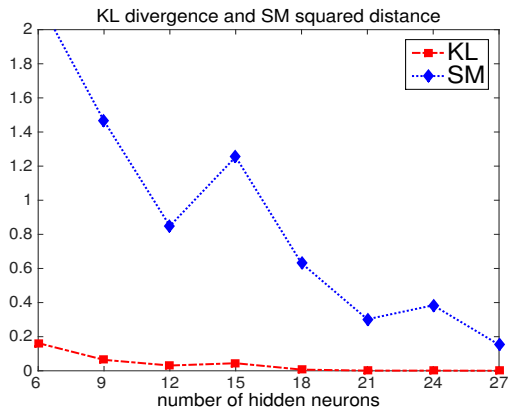
- For illustration, we consider the following beta-binomial model:

$$P(y_j|m, K) = \binom{n_j}{y_j} \frac{B(Km + y_j, K(1 - m) + n_j - y_j)}{B(Km, K(1 - m))}$$

- The following plot shows approximate posterior distributions for different numbers of hidden neurons (basis functions)

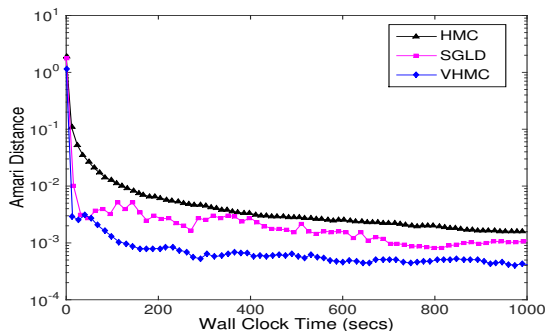


Example: beta-binomial model



Example: Independent Component Analysis

- In this example, we apply ICA to MEG data
- The following plot compares our method to HMC and SGLD (Welling and Teh, 2011) using the Amari distance (Amari et al., 1996), $d_A(\overline{W}, W_0)$, for the unmixing matrix W



Conclusion

- It is essential to develop scalable Bayesian inference methods as high-dimensional problems and complex models become commonplace in scientific studies
- Subsampling strategies have provided promising results
- However, we believe that focusing on parameter space, as opposed to data space, and exploiting its structure and regularity would lead to more reliable methods
- While methods based on surrogate functions could scale well, they might not be very effective for big data analysis
- Our variational HMC method provides a framework to bring together MCMC and Variational Bayes in order to construct robust and scalable Bayesian inference methods with both approximation accuracy and computational efficiency

Conclusion

- It is essential to develop scalable Bayesian inference methods as high-dimensional problems and complex models become commonplace in scientific studies
- Subsampling strategies have provided promising results
- However, we believe that focusing on parameter space, as opposed to data space, and exploiting its structure and regularity would lead to more reliable methods
- While methods based on surrogate functions could scale well, they might not be very effective for big data analysis
- Our variational HMC method provides a framework to bring together MCMC and Variational Bayes in order to construct robust and scalable Bayesian inference methods with both approximation accuracy and computational efficiency

Conclusion

- It is essential to develop scalable Bayesian inference methods as high-dimensional problems and complex models become commonplace in scientific studies
- Subsampling strategies have provided promising results
- However, we believe that focusing on parameter space, as opposed to data space, and exploiting its structure and regularity would lead to more reliable methods
- While methods based on surrogate functions could scale well, they might not be very effective for big data analysis
- Our variational HMC method provides a framework to bring together MCMC and Variational Bayes in order to construct robust and scalable Bayesian inference methods with both approximation accuracy and computational efficiency

Conclusion

- It is essential to develop scalable Bayesian inference methods as high-dimensional problems and complex models become commonplace in scientific studies
- Subsampling strategies have provided promising results
- However, we believe that focusing on parameter space, as opposed to data space, and exploiting its structure and regularity would lead to more reliable methods
- While methods based on surrogate functions could scale well, they might not be very effective for big data analysis
- Our variational HMC method provides a framework to bring together MCMC and Variational Bayes in order to construct robust and scalable Bayesian inference methods with both approximation accuracy and computational efficiency

Conclusion

- It is essential to develop scalable Bayesian inference methods as high-dimensional problems and complex models become commonplace in scientific studies
- Subsampling strategies have provided promising results
- However, we believe that focusing on parameter space, as opposed to data space, and exploiting its structure and regularity would lead to more reliable methods
- While methods based on surrogate functions could scale well, they might not be very effective for big data analysis
- Our variational HMC method provides a framework to bring together MCMC and Variational Bayes in order to construct robust and scalable Bayesian inference methods with both approximation accuracy and computational efficiency

Acknowledgements

- Collaborators:
 - ▶ Hongkai Zhao
 - ▶ Cheng Zhang
- Fundings:



DMS 1622490



National Institutes of Health

R01 AI107034

Thank you!